

HPC - Assignment n°2

# Durbin-Levinson algorithm

—

Salvatore Bianco, Linda Burchiellaro, Carlo Uguzzoni

# Where did we leave off?

Working on our cuda kernel, we found out:

```
sum[0][k] = r[k];  
  
for (i = 0; i <= k - 1; i++)  
    sum[i + 1][k] = sum[i][k] + r[k - i - 1] * y[i][k - 1];
```

← This could be reduced!

```
for (i = 0; i <= k - 1; i++)  
    y[i][k] = y[i][k - 1] + alpha[k] * y[k - i - 1][k - 1];
```


← Also access pattern on matrix y could be changed!

So we obtained further speed-up!


Version	Time (size = 100)	Time (size = 1000)	Time (size = 10000)	Time (size = 15000)
Sequential	0.00005	0.01169	4.41561	38.55401
Reworked sequential	0.00004	0.00267	0.29813	0.73846

# About new changes and parallelization via openmp

No speedup at all for small datasets...



...but quite better for mid-large ones!



Expected speedup should be higher. We suspect this is happening because of thread awakening and synchronization, but it could also be a matter of bandwidth.

# From base code to cuda kernel

```
int i, k;
#pragma scop
y[0][0] = r[0];
beta[0] = 1;
alpha[0] = r[0];

for (k = 1; k < _PB_N; k++)
{
    beta[k] = beta[k - 1] - alpha[k - 1] * alpha[k - 1] * beta[k - 1];
    sum[0][k] = r[k];

    for (i = 0; i <= k - 1; i++)
        sum[i + 1][k] = sum[i][k] + r[k - i - 1] * y[i][k - 1];

    alpha[k] = -sum[k][k] * beta[k];

    for (i = 0; i <= k - 1; i++)
        y[i][k] = y[i][k - 1] + alpha[k] * y[k - i - 1][k - 1];

    y[k][k] = alpha[k];
}

for (i = 0; i < _PB_N; i++)
    out[i] = y[i][_PB_N - 1];
}
```

- alpha, beta, sum to `__device__` variables
- only 2 1-d arrays for the two rows of `y` we actually work on
- only a minimal part of the old data structures to be allocated/initialized

Kernel 1 (sum reduction)

Kernel 2 (saxpy-like operation)

More in detail further on...

Deleted to spare a memcpy. Results are returned as the last row of `y` we computed.

# Data structures cuda mallocs/memcpies

```
__device__ DATA_T d_alpha, d_beta, d_sum;
```

\_\_device\_\_ variables helped removing sequential stages around the two kernels

```
// Device data structures.
DATA_T *d_r, *y_old, *y_new;

// Device mallocs.
gpuErrchk(cudaMalloc((void **)&d_r, sizeof(DATA_T) * N));
gpuErrchk(cudaMalloc((void **)&y_old, sizeof(DATA_T) * N));
gpuErrchk(cudaMalloc((void **)&y_new, sizeof(DATA_T) * N));
```

Just 3 linear arrays (size n) allocated

```
// Memcopies.
// Device's array r.
gpuErrchk(cudaMemcpy(d_r, h_r, sizeof(DATA_T) * N, cudaMemcpyHostToDevice));
// y_old[0] = r[0].
gpuErrchk(cudaMemcpy(y_old, d_r, sizeof(DATA_T), cudaMemcpyDeviceToDevice));
DATA_T alpha;
DATA_T beta = 1;
// alpha = r[0].
gpuErrchk(cudaMemcpy(&alpha, d_r, sizeof(DATA_T), cudaMemcpyDeviceToHost));
gpuErrchk(cudaMemcpyToSymbol(d_alpha, &alpha, sizeof(DATA_T)));
// beta = 1.
gpuErrchk(cudaMemcpyToSymbol(d_beta, &beta, sizeof(DATA_T)));

// Function kernel durbin call (device).
kernel_durbin_device(y_old, y_new, d_r);

// out = y_new dell'ultima iterazione di durbin (viene swappato -> quindi y_old)
gpuErrchk(cudaMemcpy(d_out, y_new, sizeof(DATA_T) * N, cudaMemcpyDeviceToHost));
```

- Minimum data motion from host to device
- \_\_device\_\_ variables initialization (where needed)

# Kernel 1 - sum reduction

```
// Device kernel.
// Variabili dislocate su device.
__device__ DATA_T d_alpha, d_beta, d_sum;

// Primo kernel -> calcolo del nuovo beta + calcolo delle somme parziali e successiva reduction su sum.
__global__ void first_kernel(DATA_T *__restrict__ y, DATA_T *__restrict__ r, int k)
{
    __shared__ DATA_T partialSum[BLOCK_SIZE];

    // Coordinate del thread.
    int tid = threadIdx.x;
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    // Calcolo del valore base di sum.
    d_sum = r[k];

    // Calcolo del nuovo beta.
    DATA_T beta = d_beta;
    d_beta = beta - d_alpha * d_alpha * beta;

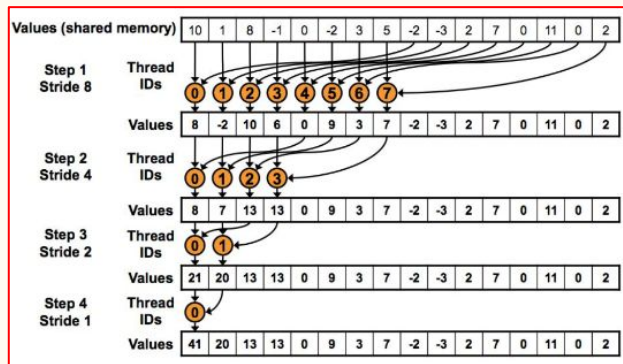
    // Caricamento delle somme parziali in memoria condivisa.
    if (i < k)
        partialSum[tid] = r[k - i - 1] * y[i];
    else
        partialSum[tid] = 0;

    __syncthreads();

    // Riduzione. Ciascun blocco porta la propria somma parziale in partialSum[0].
    for (int stride = blockDim.x / 2; stride > 0; stride >>= 1)
    {
        if (tid < stride)
            partialSum[tid] += partialSum[tid + stride];
        __syncthreads();
    }

    // Il thread con tid 0 di ciascun blocco aggiorna il valore globale nel device con una atomicAdd.
    if (tid == 0)
        atomicAdd(&d_sum, partialSum[0]);
}
```

- Use of shared memory for best bandwidth
- Partial sums are stored in an array of size BLOCK\_SIZE
- for loop operates a first sum reduction for each block on array index 0
- Every thread with tid = 0 adds it's block partial sum to \_\_device\_\_ d\_sum with an atomicAdd



# Kernel 2 - saxpy

```
// Secondo kernel -> calcolo del nuovo alpha + calcolo del nuovo y in stile saxpy.
__global__ void second_kernel(DATA_T *__restrict__ y_old, DATA_T *__restrict__ y_new, int k)
{
    // Coordinate del thread.
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    // Calcolo del nuovo alpha.
    d_alpha = -d_sum * d_beta;

    if (i < k)
        y_new[i] = y_old[i] + d_alpha * y_old[k - i - 1];

    y_new[k] = d_alpha;
}
```

Basic saxpy-like kernel

```
// Funzione chiamante dei kernel. Replica il kernel di Durbin.
void kernel_durbin_device(
    DATA_T *__restrict__ y_old,
    DATA_T *__restrict__ y_new,
    DATA_T *__restrict__ d_r)
{
    int k;
    // int GRID_SIZE = (N + BLOCK_SIZE - 1) / BLOCK_SIZE;
    int GRID_SIZE;

    for (k = 1; k < N; k++)
    {
        GRID_SIZE = (k + BLOCK_SIZE - 1) / BLOCK_SIZE;

        // Calcolo del nuovo beta e di sum.
        first_kernel<<<GRID_SIZE, BLOCK_SIZE>>>(y_old, d_r, k);

        // Calcolo del nuovo alpha e del nuovo y.
        second_kernel<<<GRID_SIZE, BLOCK_SIZE>>>(y_old, y_new, k);

        // Scambio degli y.
        swapPointers(y_old, y_new);
    }
}
```

Grid size is determined by variable k for both kernel calls, to avoid creating non-working blocks

# Results (overview)


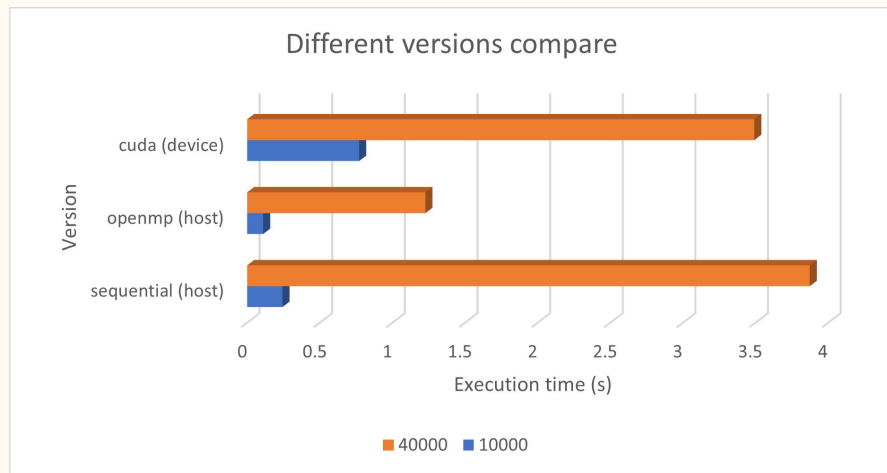
```
Durbin (Host) : 0.243 sec 8242.6 GFLOPS
Durbin (GPU): 0.762 sec 2624.8 GFLOPS
==13349== Profiling application: ./durbin_vfinal.exe
==13349== Profiling result:
   Type  Time(%)   Time     Calls   Avg      Min      Max  Name
GPU activities: 72.36% 120.75ms 9999 12.075us 2.3430us 22.868us first_kernel(float*, float*, int)
                27.63% 46.103ms 9999 4.6100us 1.0930us 7.9180us second_kernel(float*, float*, int)
                0.00% 7.5530us 3 2.5170us 209ns 7.0840us [CUDA memcpy HtoD]
                0.00% 4.8970us 2 2.4480us 834ns 4.0630us [CUDA memcpy DtoH]
                0.00% 2.3440us 1 2.3440us 2.3440us 2.3440us [CUDA memcpy DtoD]
API calls: 69.27% 725.53ms 19998 36.280us 34.167us 895.07us cudaLaunchKernel
            30.63% 320.76ms 3 106.92ms 18.646us 320.72ms cudaMalloc
            0.06% 588.60us 4 147.15us 97.919us 173.13us cudaMemcpy
            0.02% 249.07us 3 83.022us 21.667us 180.00us cudaFree
            0.01% 122.97us 97 1.2670us 677ns 31.772us cuDeviceGetAttribute
            0.01% 80.783us 2 40.391us 35.105us 45.678us cudaMemcpyToSymbol
            0.00% 10.052us 1 10.052us 10.052us 10.052us cuDeviceTotalMem
            0.00% 5.9880us 3 1.9960us 1.5100us 2.9160us cuDeviceGetCount
            0.00% 2.5000us 1 2.5000us 2.5000us 2.5000us cuDeviceGetName
            0.00% 2.4470us 2 1.2230us 989ns 1.4580us cuDeviceGet
            0.00% 937ns 1 937ns 937ns 937ns cuDeviceGetUuid
```

Generic profile test executed on a 10k-sized dataset

- almost 100% computation on GPU happens inside the two kernels
  - Also the kernels sum up to ~160ms (which is faster than openmp parallelized version)
- almost 100% of API calls happens in:
  - cudaMalloc (unavoidable)
  - cudaLaunchKernel: 2 calls per iteration in the outermost for loop, which carries dependencies



# Results (charts)



Dataset sizes

# Possible further improvements

- First kernel could be improved:
  - `atomicAdd` is sub-optimal. Could write a new implementation which doesn't use it, thus 100% based on block reductions
  - By also applying reductions on single warps instead of blocks (they need no synchronization)
- Second kernel could be improved:
  - Memory reads happen twice on the same array `y` elements (`y_old[i]` and `y_old[k - i - 1]`). Surely there's some way to halven memory accesses

## Conclusions

With all summed up, this cuda version is still faster than the original version of the kernel, also it tends to get faster than the rewritten kernel on very large datasets. Yet the `openmp` version is always faster, due to the lack of overhead.